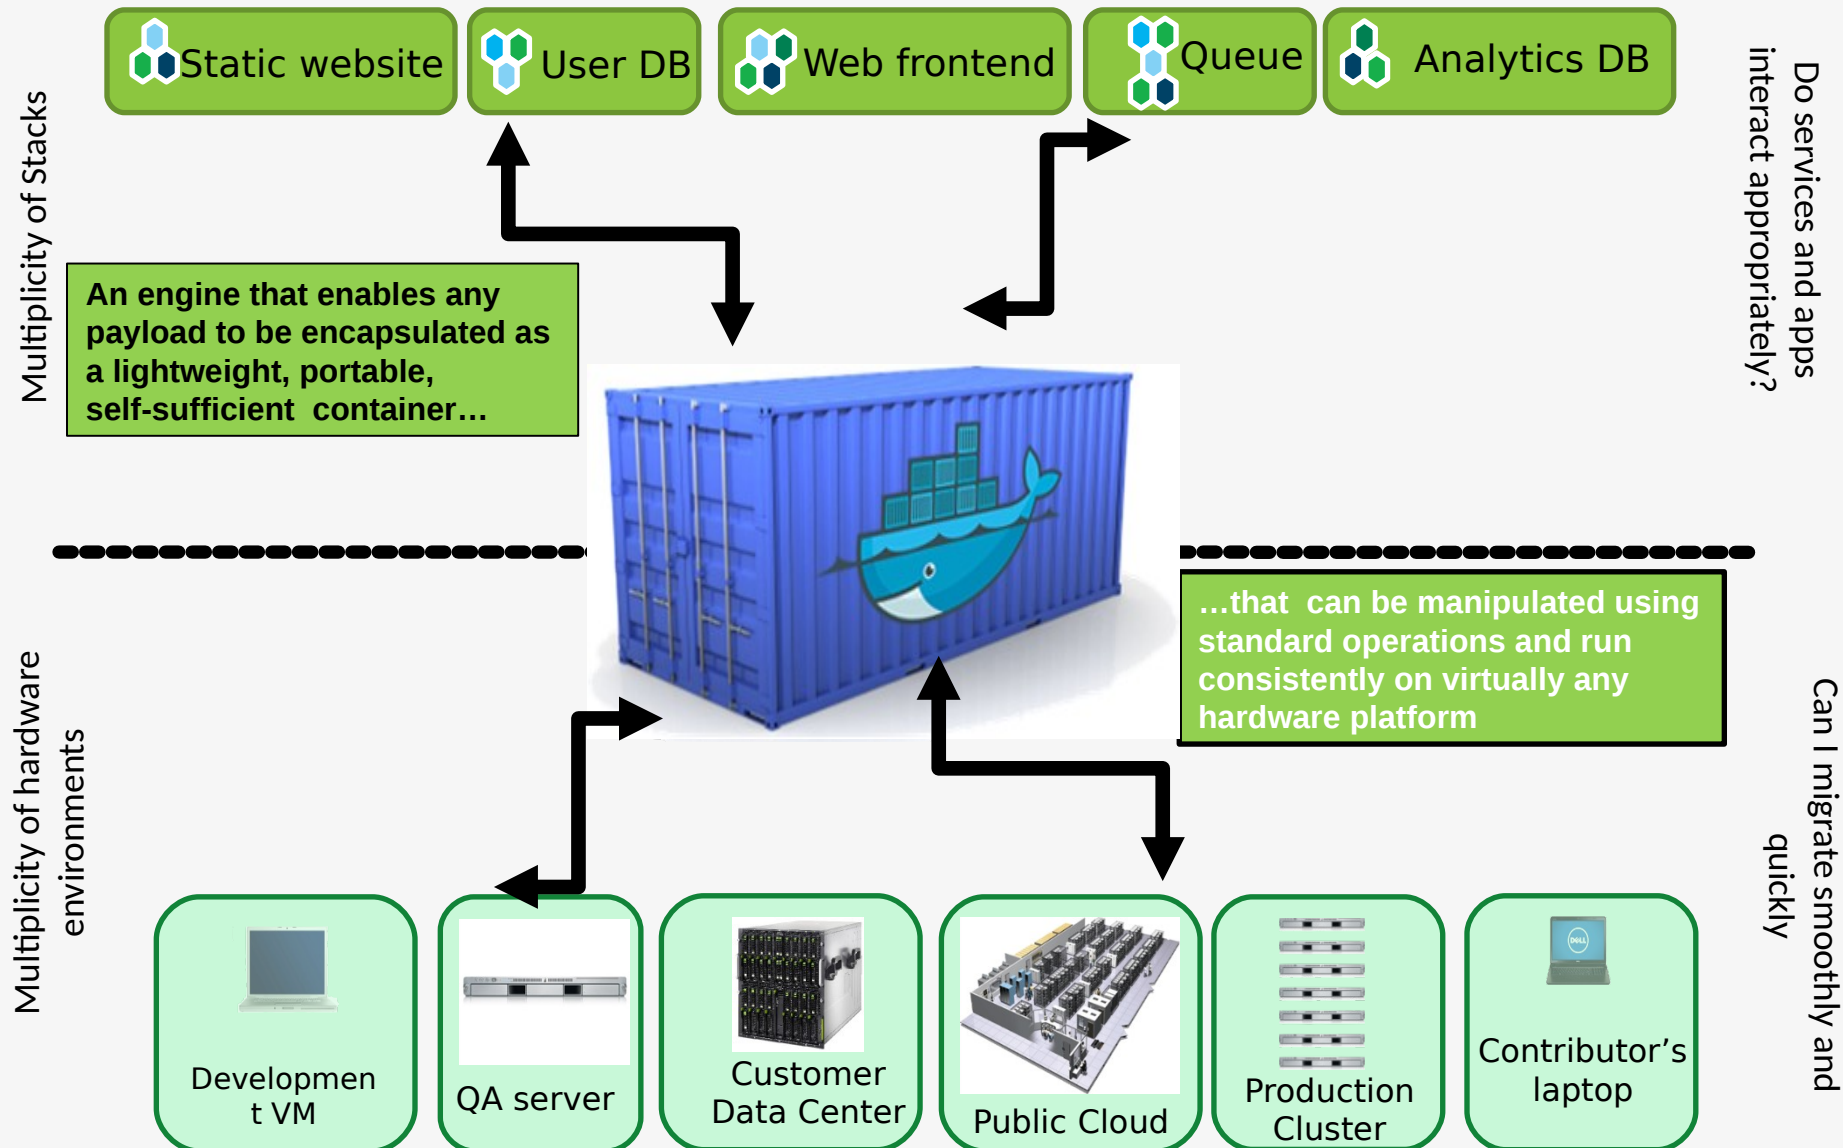


# Java and Docker

Dinakar Guniguntala, Architect, Java Cloud Optimization

# Docker is a shipping container system for code




# Docker Basics



## Image

- A read-only snapshot of a container stored in Docker Hub to be used as a template for building containers



## Container

- The standard unit in which the application service resides or transported

SaaS

Enterprise

## Docker Hub

- Available in SaaS or Enterprise to deploy anywhere you choose
- Stores, distributes and shares container images



## Docker Engine

- A program that creates, ships and runs application containers
- Runs on any physical and virtual machine or server locally, in private or public cloud
- Client communicates with Engine to execute commands

# Docker Basics Contd



## Isolation

- Docker uses several Linux kernel features such as namespaces (ipc, mount, pid, network and user), Apparmor and SELinux profiles and control groups (cgroups), chroot on steroids to achieve Isolation



## Lightweight

- Containers are “light” users of system resources, much smaller than VMs, startup time is less than a second and have better performance



## User-friendly

- Developers build with ease and ship higher-quality applications
- Sysadmins deploy workload based on business priorities and policies

# Dockerfile basics

- Layers – aufs (or other unionFS)
- `$ docker history`
- Docker image size
  - push times
  - provisioning times
  - deploy times
  - storage
  - security
- Alpine Linux, SFJ images

# Java Container Startup Time

- Shared Class Cache
  - -Xshareclasses:name=appcache
  - `java -Xshareclasses:name=liberty,cacheDir=/opt/ibm/wlp/output/.classCache,printallstats`
- Liberty startup down to 3.x seconds from 4.x seconds
  - 20 – 30% reduction in startup time with SCC
- Run application in typical workload scenarios with SCC turned on.
  - Gather classes loaded (using “printallstats”)
  - Bake SCC into base Docker Image.

# CPU and Memory

- cgroups
  - \$ systemd-cgls
  - \$ systemd-cgtop
  - cpu.shares
- cpusets
  - cpus
  - mems
    - /sys/fs/cgroup/memory/memory.limit\_in\_bytes
    - /sys/fs/cgroup/memory/docker/memory.soft\_limit\_in\_bytes
  - java -verbose:sizes -version
- Docker options
  - cpu-percent int CPU percent (Windows only)
  - cpu-period int Limit CPU CFS period
  - cpu-quota int Limit CPU CFS quota
  - c, --cpu-shares int CPU shares (relative weight)
  - cpuset-cpus string CPUs in which to allow execution (0-3, 0,1)
  - cpuset-mems string MEMs in which to allow execution (0-3, 0,1)
  - memory string Memory limit
  - memory-reservation string Memory soft limit
  - memory-swap string Swap limit equal to memory plus swap
    - '-1' to enable unlimited swap
  - memory-swappiness int Tune container memory swappiness (0 to 100)  
(default -1)

# The Idle problem

- 30% servers are sitting idle in the cloud\*
- Charging model on most clouds based on memory usage (GB/hr on Bluemix).
- Historically poor Java idle behavior causes CPU burn.
  - Starves other JVM instances.
  - Increased costs for CPU usage.
- Java is currently inefficient with heap usage on Idle.
  - No reduction in memory usage even when on long periods of idle.
  - Big data workloads cause heap expansion but don't contract on Idle.
- Might hurt application performance on Active after a long period of Idle as heap cleanup may be needed.
- First step is to be able to measure precise JVM CPU and Memory usage.

\* <https://www.forbes.com/sites/benkepes/2015/06/03/30-of-servers-are-sitting-comatose-according-to-research/#39b31e8d59c7>



# JvmCpuMonitorInfo

**JvmCpuMonitorInfo()**

Creates a new JvmCpuMonitorInfo instance.

## Method Summary

All Methods	Static Methods	Instance Methods	Concrete Methods
-------------	----------------	------------------	------------------

Modifier and Type	Method and Description
boolean	<b>equals</b> (java.lang.Object obj) Compares this object with the specified object for equality.
static JvmCpuMonitorInfo	<b>from</b> (javax.management.openmbean.CompositeData cd) Receives a CompositeData representing a JvmCpuMonitorInfo object and attempts to return the root JvmCpuMonitorInfo instance.
long	<b>getApplicationCpuTime</b> () This method returns the total CPU usage for all application threads.
long[]	<b>getApplicationUserCpuTime</b> () This method returns an array of CPU usage for all user defined thread categories.
long	<b>getGcCpuTime</b> () This method returns the total CPU usage of all GC threads.
long	<b>getJitCpuTime</b> () This method returns the total CPU usage of all JIT Threads.
long	<b>getResourceMonitorCpuTime</b> () This method returns the total CPU usage for all threads of the "Resource-Monitor" category.
long	<b>getSystemJvmCpuTime</b> () This method returns the total CPU usage of the "System-JVM" category, which includes GC, JIT and other JVM daemon threads.
long	<b>getTimestamp</b> () This method returns the last sampling time stamp.
int	<b>hashCode</b> () Returns the hash code value for this object.
java.lang.String	<b>toString</b> () Text description of this JvmCpuMonitorInfo object.

## -Xtune:virtualized

- Reduces JVM CPU consumption when Idle.
- Needs a large shared class cache to maintain peak performance.
- AoT space in the Shared Class Cache (SCC) must not be capped.

# Logging and Monitoring

- Logging
  - Docker Volume Container (DVC)
    - --volumes-from
    - `docker create -v /logdata --name logdvc websphere-liberty /bin/true`
- Monitoring
  - `$ docker top`
  - `$ docker stats`
  - `$ docker inspect`

**QUESTIONS ?**