

# Java Concurrency

Harshit Bangar

Credit - Douglas Hawkins

# Quiz - Identify thread safe APIs

- long height = 64L
- x += 1
- Point point = new Point();

# Quiz - Identify thread safe APIs

- long height = 64L
- x += 1
- Point point = new Point();

# Quiz - Identify thread safe APIs

- long height = 64L
- x += 1
- Point point = new Point();

# Quiz - Identify thread safe APIs

- long height = 64L
- x += 1
- Point point = new Point();

# Content

# Content

- Unit testing concurrent code.

# Content

- Unit testing concurrent code.
- Atomicity

# Content

- Unit testing concurrent code.
- Atomicity
- Visibility

# Content

- Unit testing concurrent code.
- Atomicity
- Visibility
- Thread Safety

# Content

- Unit testing concurrent code.
- Atomicity
- Visibility
- Thread Safety
- Designing a thread safe class

# Content

- Unit testing concurrent code.
- Atomicity
- Visibility
- Thread Safety
- Designing a thread safe class
- Java Memory Model

# Unit testing

# Unit testing

1. The biggest challenge with testing concurrent code is unpredictability. Unlike code which is sequential and deterministic, concurrency can yield completely random results

# Unit testing

1. The biggest challenge with testing concurrent code is unpredictability. Unlike code which is sequential and deterministic, concurrency can yield completely random results
2. Unit tests sometimes introduces sequentiality by accident

# Unit testing

# Unit testing

- Junit

# Unit testing

- Junit
- TestNg

# Unit testing

- Junit
- TestNg
- ThreadWeaver

# Unit testing

- Junit
- TestNg
- ThreadWeaver
- JCStress

# Unit testing - Junit & TestNg



# Unit testing - Junit & TestNg

- Junit doesn't support concurrent testing.

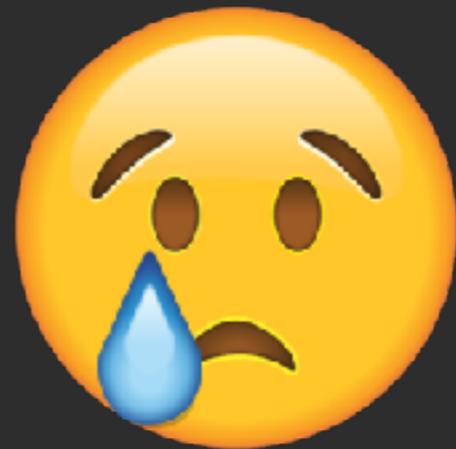


# Unit testing - Junit & TestNg

- Junit doesn't support concurrent testing.
- TestNg supports testing but it is broken since it is unable to reproduce the consistent behaviour.

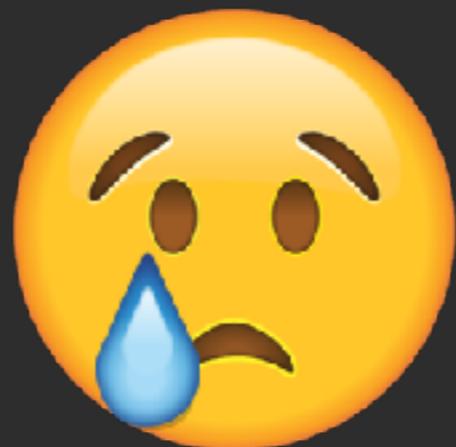


# Unit testing - Thread Weaver



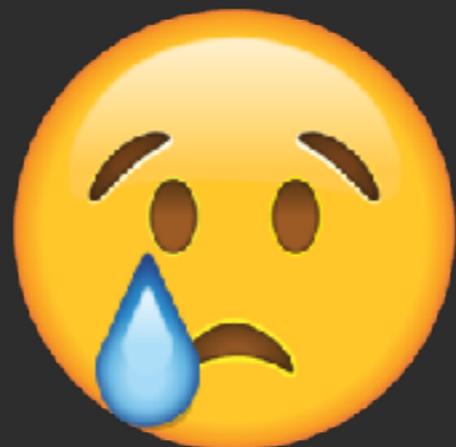
# Unit testing - Thread Weaver

- Works by adding conditional breakpoints in the code (by inserting locks via byte code manipulation)



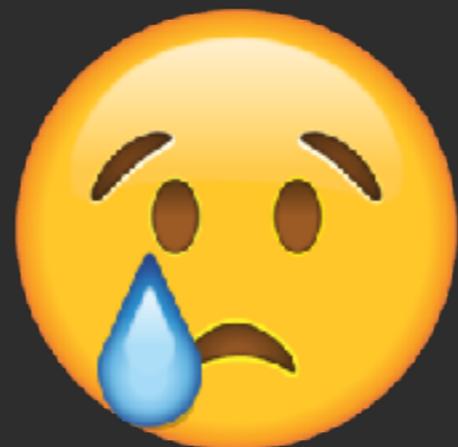
# Unit testing - Thread Weaver

- Works by adding conditional breakpoints in the code (by inserting locks via byte code manipulation)
- Theoretically should catch every race condition.



# Unit testing - Thread Weaver

- Works by adding conditional breakpoints in the code (by inserting locks via byte code manipulation)
- Theoretically should catch every race condition.
- It prevent JMM re-ordering (accidentally) due to locks - details.



# Unit testing - JCStress

# Unit testing - JCStress

- Stress testing by simulating concurrency via multiple actors

# Unit testing - JCStress

- Stress testing by simulating concurrency via multiple actors
- Not a unit test, result might vary from run to run.

# Unit testing - JCStress

- Stress testing by simulating concurrency via multiple actors
- Not a unit test, result might vary from run to run.
- Result might depend on the processor architecture.

# Unit testing - JCStress

- Stress testing by simulating concurrency via multiple actors
- Not a unit test, result might vary from run to run.
- Result might depend on the processor architecture.
- Examples

# ATOMICITY

## WHAT OPERATIONS ARE INDIVISIBLE?

# SUCCINCT != ATOMIC

sharedX = 2L;

thread 1

sharedX = 2L

set\_hi sharedX, 0000 0000

set\_lo sharedX, 0000 0002

thread 2

sharedX = -1L

set\_hi sharedX, ffff ffff  
set\_lo sharedX, ffff ffff

# SUCCINCT != ATOMIC

```
sharedX += 1;
```

```
localX = sharedX;  
localX = localX + 1;  
sharedX = localX;
```

# SUCCINCT != ATOMIC

```
Point sharedPoint = new Point(x, y);

local1 = calloc(sizeof(Point));
local1.<init>(x, y);
Object.<init>();
this.x = x;
this.y = y;
sharedPoint = local1;
```

# VISIBILITY

## WHAT CAN OTHER THREADS SEE?

**NO GARBAGE VALUES**  
**NO OUT OF “THIN AIR” VALUES**

WILL SEE A ZERO-ISH VALUE

OR

AN ASSIGNED VALUE

---

**NO \*COMPLETELY\* GARBAGE VALUES**

TWO SEPARATE HALVES OF LONG OR DOUBLE

BUT

BOTH HALVES ARE ZERO OR ASSIGNED

# ONLY ASSIGNED VALUES FOR FINALS\*

```
sharedPoint = new Point(x, y);
```

```
class Point {  
    final int x, y;  
  
    public Point(int x, int y) {  
        this.x = x;  
        this.y = y;  
    }  
}
```

JLS8 § 17.5 pp 652

\*After constructor return

# NOT \*ALL\* VALUES ARE STORED

~~sharedX = 20;~~ ← Dead Store

sharedX = 40;

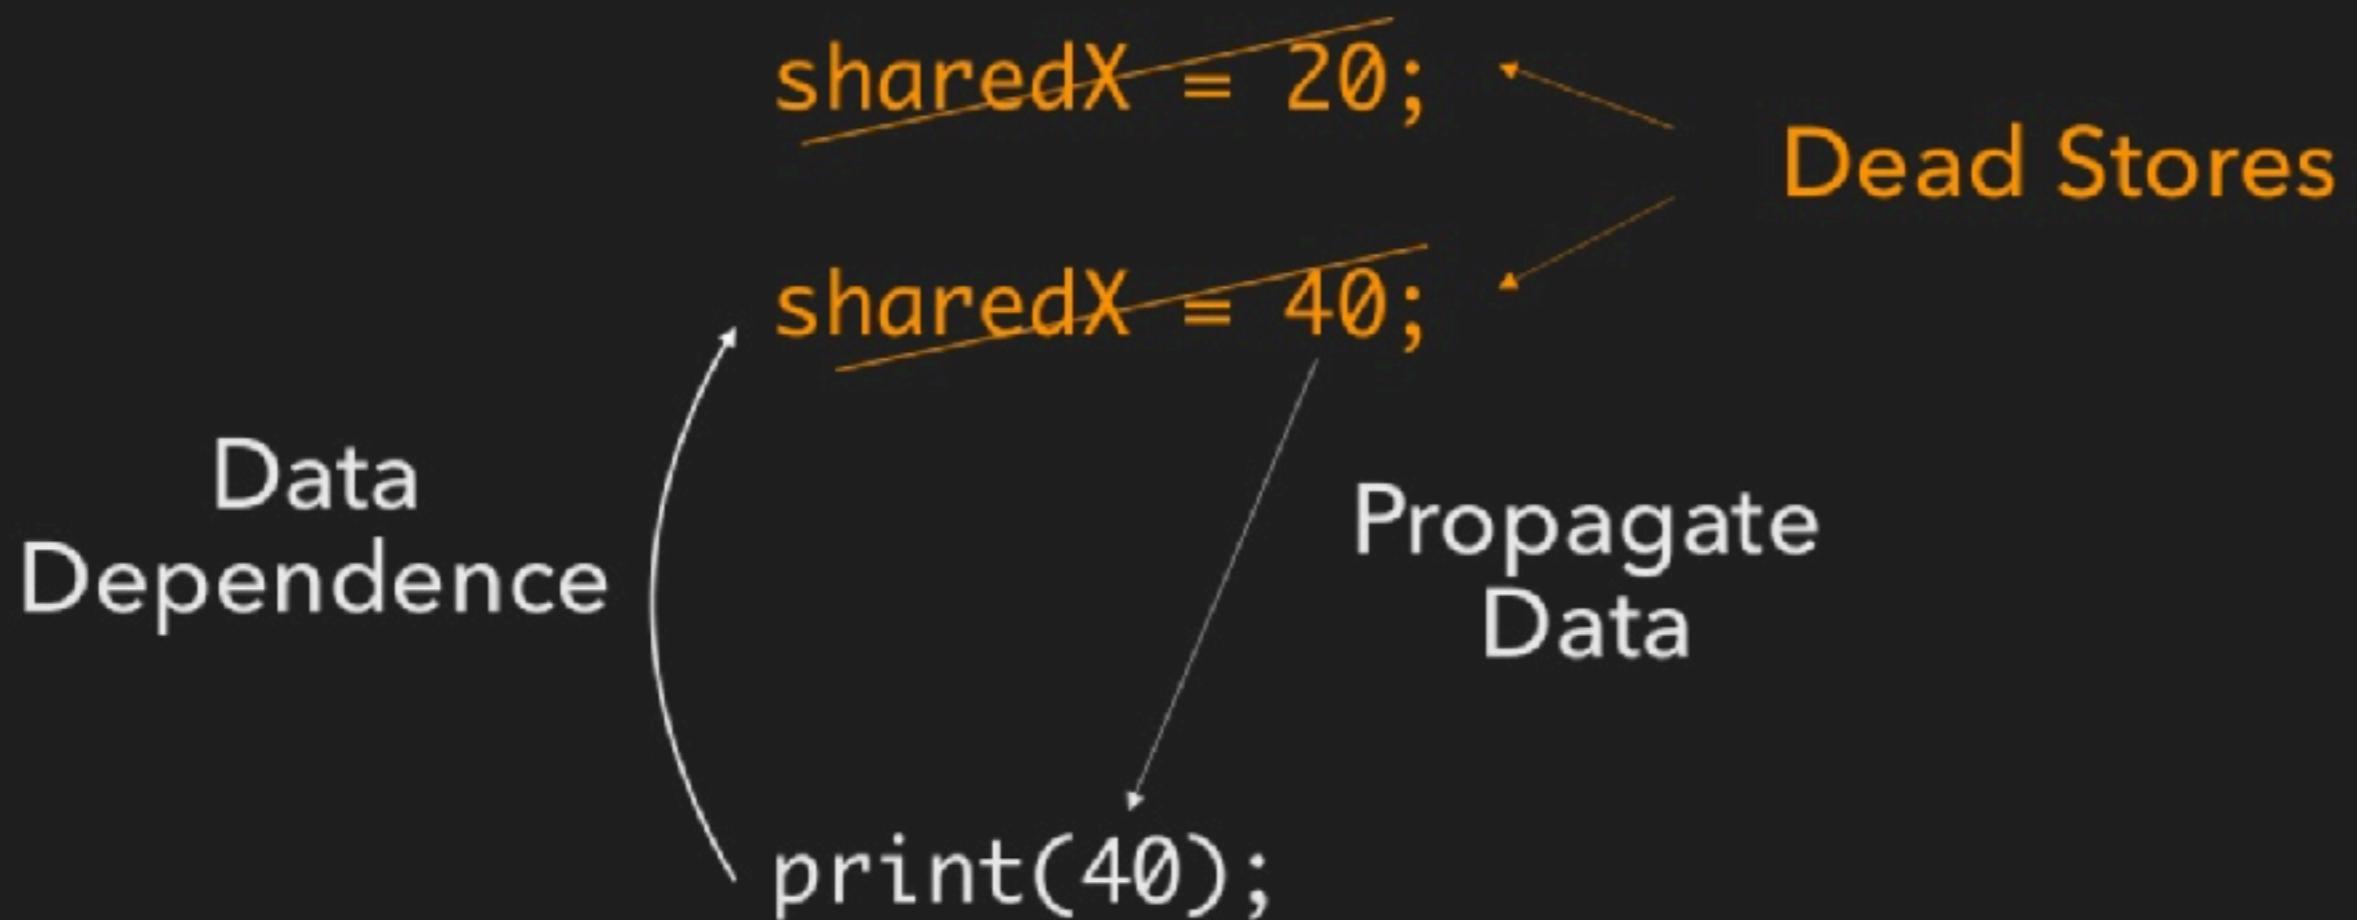
print(sharedX);

# DATA DEPENDENCE

sharedX = 20; ← Dead Store

Data  
Dependence ↗ sharedX = 40;  
print(sharedX);

# DATA DEPENDENCE



## NOT \*ALL\* VALUES ARE READ

```
xSquared = sharedX * sharedX
```

Can xSquared be 30?

```
local1 = sharedX
```

```
local2 = sharedX
```

```
xSquared = local1 * local2
```

# NOT \*ALL\* VALUES ARE READ

Can xSquared be 30?

```
local1 = sharedX  
local2 = sharedX  
xSquared = local1 * local2
```

	thread 1	thread 2
	local1 = sharedX; (5) local2 = sharedX; (6) xSquared = (30) local1 * local2;	sharedX = 5; sharedX = 6;

# ORDERING

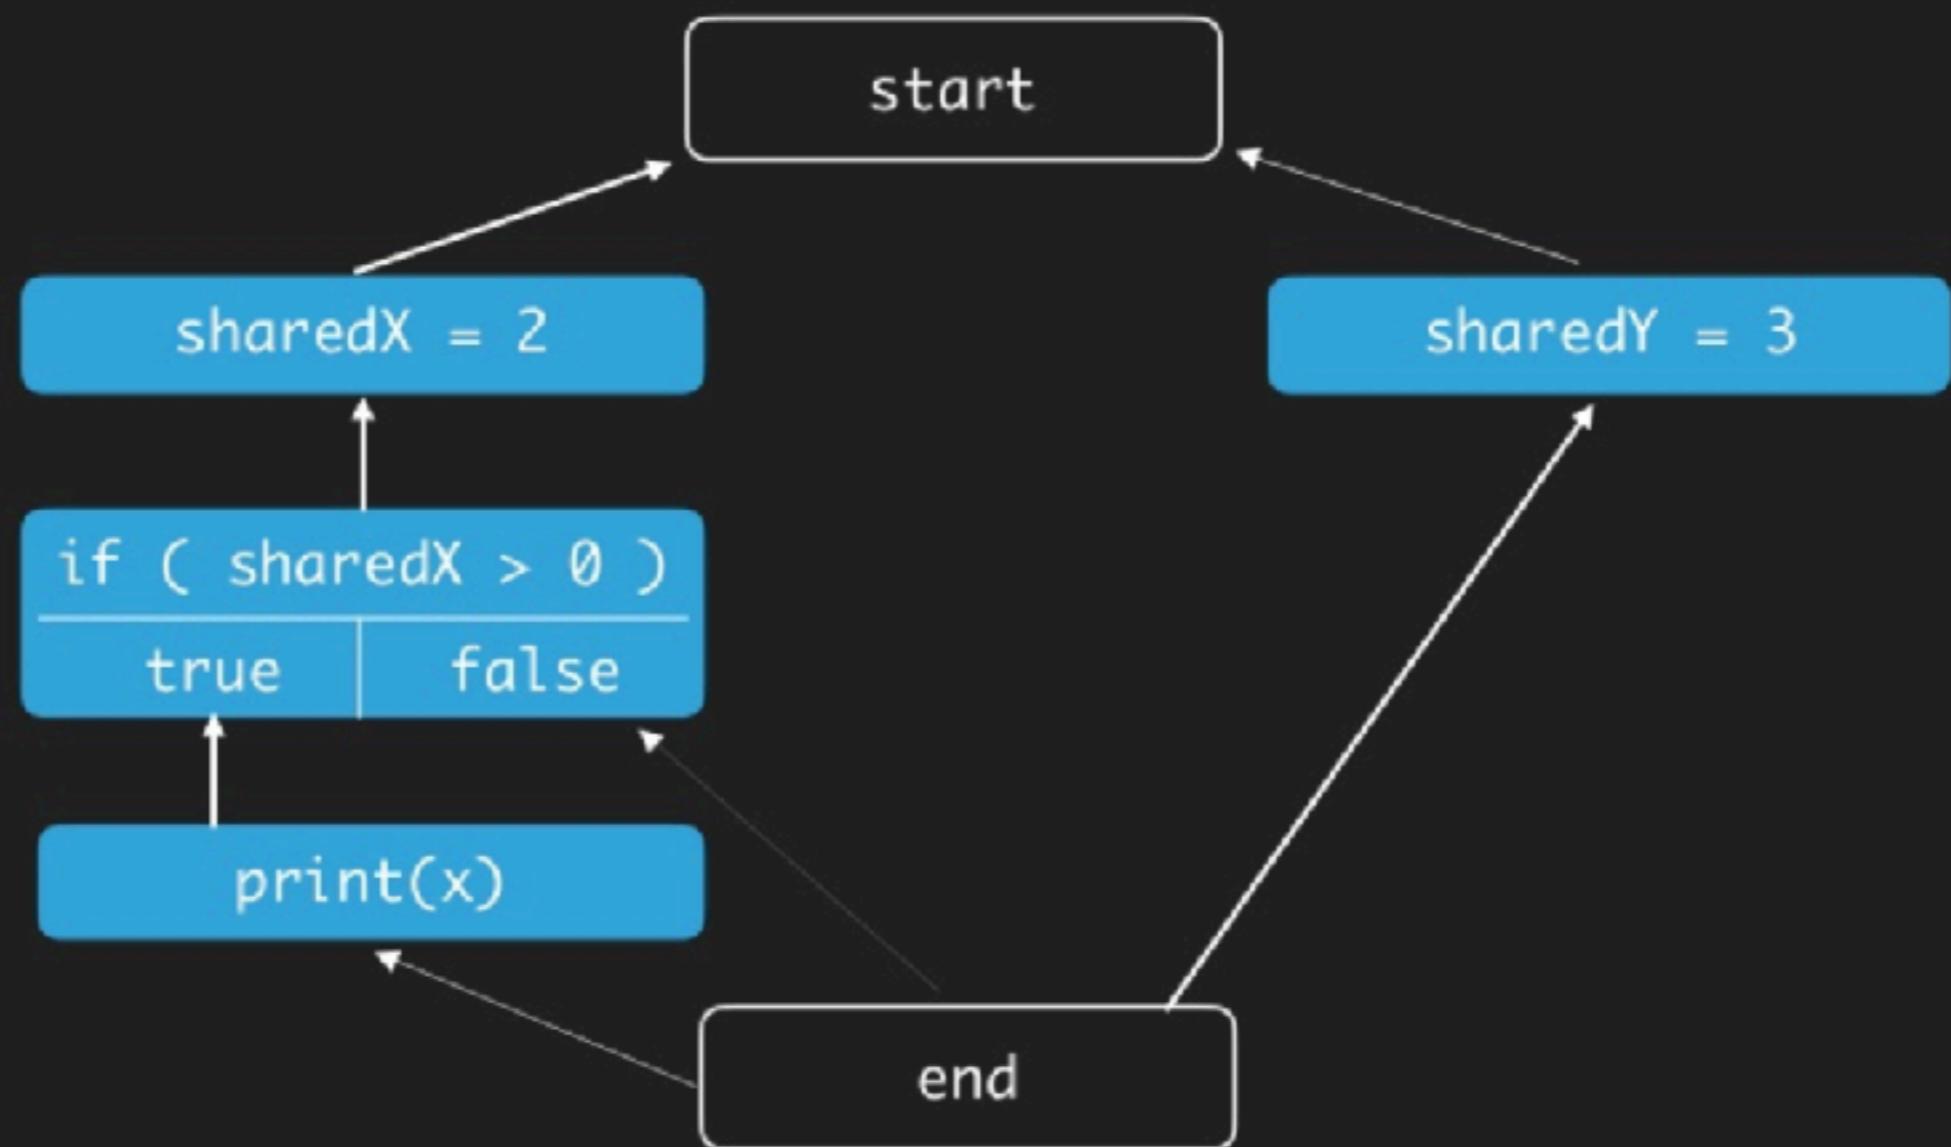
## WHAT ORDER DOES IT RUN IN?

# “PROGRAM ORDER”

## DATA + CONTROL DEPENDENCE

Data                      sharedX = 2              ←  
Control              ↗ sharedY = 3  
                    ↗ if ( sharedX > 0 ) {  
                            print(sharedX);  
                    }              ↘ Data

# DATA + CONTROL DEPENDENCE



# STORE SOONER

```
sharedFoo = new Foo(...);
```



```
    local1 = calloc(sizeof(Point));  
    local1.<init>(...);  
    ...  
    sharedPoint = local1;
```

The diagram illustrates the concept of "STORE SOONER". It shows a sequence of C-like code. A curved arrow originates from the word "Data" on the left and points to the assignment operator (=) in the line "sharedPoint = local1;". Another "Data" label is placed near the first assignment operator ("local1 ="). Ellipses between the assignments indicate additional code.

# FREE TO REORDER

```
sharedFoo = new Foo(...);
```

The diagram illustrates the memory layout and potential reordering of variables. It shows two sections of memory, each labeled "Data". The first section contains the assignment `local1 = calloc(sizeof(Point));`. The second section contains the assignment `sharedPoint = local1;` and the initialization `local1.<init>(...);`. A curved arrow points from the first "Data" section to the second, indicating the flow of data. An orange curved arrow points from the second "Data" section back to the first, with the text "Reordered!" written next to it, indicating that the assignments could be reordered by the compiler.

```
local1 = calloc(sizeof(Point));  
sharedPoint = local1;  
local1.<init>(...);
```

Data

Data

Reordered!

# \*BROKEN\* DOUBLE CHECKED LOCKING

```
static Singleton instance() {  
    if ( sharedInstance == null ) {  
        synchronized ( Singleton.class ) {  
            if ( sharedInstance == null ) {  
                sharedInstance = new Singleton();  
            }  
        }  
    }  
    return sharedInstance;  
}
```

```
static Singleton instance() {  
    if ( sharedInstance == null ) {  
        synchronized ( Singleton.class ) {  
            if ( sharedInstance == null ) {  
                local = calloc(sizeof(Singleton));  
                sharedInstance = local; ↓  
                local.<init>();  
            }  
        }  
    }  
    return sharedInstance;  
}
```

sharedInstance is non-null,  
but constructor hasn't run.

# PRODUCER

```
...  
sharedData = ...;  
sharedDone = true; ➔ Reorder!
```



```
...  
sharedDone = true;  
sharedData = ...;
```

# CONSUMER

```
while ( !sharedDone );  
print(sharedData);
```

**AGAIN BLAME  
HARDWARE.**

	RISC / ARM	CISC / x86
Loads After Loads	YES	YES
Loads After Stores	YES	YES
Stores After Stores	YES	NO
Stores After Loads	YES	YES

# TO STOP REORDERING, USE A \*FENCE\*

## PRODUCER

```
...  
sharedData = ...;  
--- store_store_fence(); ---  
sharedDone = true;
```

## CONSUMER

```
while ( !sharedDone ) {  
--- load_load_fence(); ---  
}  
print(sharedData);
```

BLAME JAVA,  
TOO.

	RISC / ARM	CISC / x86	Java
Loads After Loads	YES	YES	YES
Loads After Stores	YES	YES	YES
Stores After Stores	YES	NO	YES
Stores After Loads	YES	YES	YES

**HOW DO YOU TELL  
JAVA TO NOT REORDER?**

# SYNCHRONIZATION ACTIONS

VOLATILE  
SYNCHRONIZED  
FINAL / FREEZE

~~UNSAFE~~  
ATOMICS  
VAR HANDLES

# FINAL & FREEZE

```
Point sharedPoint = new Point(x, y);
```

```
local1 = calloc(sizeof(Point));
```

```
local1.<init>(x, y);
```

```
Object.<init>();
```

```
this.x = x;
```

```
this.y = y;
```

```
----- freeze(); -----
```

```
sharedPoint = local1;
```

# FINAL & FREEZE

```
Point sharedPoint = new Point(x, y);
```

```
    local1 = calloc(sizeof(Point));
```

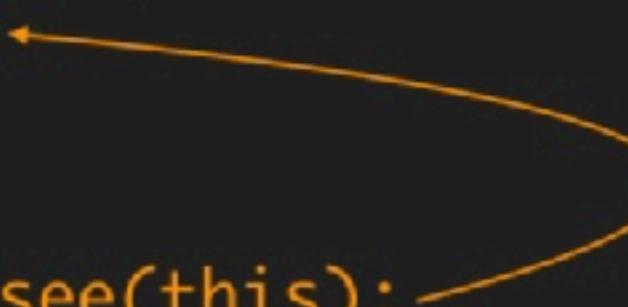
```
    local1.<init>(x, y);
```

```
    Object.<init>();
```

```
    this.x = x;
```

```
    this.y = y;
```

```
    OtherThread.see(this);
```



Free to Reorder!

```
----- freeze(); -----
```

```
sharedPoint = local1;
```

# Final & Freeze

```
class FinalWrapperFactory {  
    private FinalWrapper w;  
  
    public Singleton get() {  
        if (w == null) { // check 1  
            synchronized(this) {  
                if (w == null) { // check2  
                    w = new FinalWrapper(new Singleton());  
                }  
            }  
        }  
        return w.instance;  
    }  
  
    private static class FinalWrapper {  
        public final Singleton instance;  
        public FinalWrapper(Singleton instance) {  
            this.instance = instance;  
        }  
    }  
}
```

# Final & Freeze

```
class FinalWrapperFactory {
    private FinalWrapper w;

    public Singleton get() {
        if (w == null) { // check 1
            synchronized(this) {
                if (w == null) { // check2
                    w = new FinalWrapper(new Singleton());
                }
            }
        }
        return w.instance;
    }

    private static class FinalWrapper {
        public final Singleton instance;
        public FinalWrapper(Singleton instance) {
            this.instance = instance;
        }
    }
}
```

# Final & Freeze

```
class FinalWrapperFactory {  
    private FinalWrapper w;  
  
    public Singleton get() {  
        if (w == null) { // check 1  
            synchronized(this) {  
                if (w == null) { // check2  
                    w = new FinalWrapper(new Singleton());  
                }  
            }  
        }  
        return w.instance;  
    }  
  
    private static class FinalWrapper {  
        public final Singleton instance;  
        public FinalWrapper(Singleton instance) {  
            this.instance = instance;  
        }  
    }  
}
```

# Final & Freeze

```
class FinalWrapperFactory {  
    private FinalWrapper w;  
  
    public Singleton get() {  
        if (w == null) { // check 1  
            synchronized(this) {  
                if (w == null) { // check2  
                    w = new FinalWrapper(new Singleton());  
                }  
            }  
        }  
        return w.instance;  
    }  
  
    private static class FinalWrapper {  
        public final Singleton instance;  
        public FinalWrapper(Singleton instance) {  
            this.instance = instance;  
        }  
    }  
}
```

# VOLATILE

## PRODUCER

```
...  
volatileData = ...;  
-----  
volatileDone = true;
```

## CONSUMER

```
while ( !volatileDone ) {  
    -----  
}  
print(volatileData);
```

## ALSO FIXES DOUBLE CHECKED LOCKING

```
static Singleton instance() {
    if ( sharedInstance == null ) {
        synchronized ( Singleton.class ) {
            if ( sharedInstance == null ) {
                local = calloc(sizeof(Singleton));
                local.<init>();
                -----
                sharedInstance = local;
            }
        }
    }
    return sharedInstance;
}
```

# RIGHT WAY TO DO LAZY SINGLETON IN JAVA

## CLASS INITIALIZATION IS ALREADY LAZY.

```
static Singleton instance() {  
    return Holder.INSTANCE;  
}  
  
static class Holder {  
    static final class Singleton INSTANCE = new Singleton();  
}
```

# Mutable state

There are three ways to fix it:

# Mutable state

There are three ways to fix it:

- Use synchronisation whenever accessing the state variable.

# Mutable state

There are three ways to fix it:

- Use synchronisation whenever accessing the state variable.
- Don't share the state variable across threads;

# Mutable state

There are three ways to fix it:

- Use synchronisation whenever accessing the state variable.
- Don't share the state variable across threads;
- Make the state variable immutable; or

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        encodeIntoResponse(resp, factors);  
    }  
}
```

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        encodeIntoResponse(resp, factors);  
    }  
}
```



# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    private long count = 0;  
  
    public long getCount() {  
        return count;  
    }  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        ++count;  
        encodeIntoResponse(resp, factors);  
    }  
}
```

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    private long count = 0;  
  
    public long getCount() {  
        return count;  
    }  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        ++count;  
        encodeIntoResponse(resp, factors);  
    }  
}
```



# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    private final AtomicLong count = new AtomicLong(0);  
  
    public long getCount() {  
        return count;  
    }  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        count.incrementAndGet();  
        encodeIntoResponse(resp, factors);  
    }  
}
```



# COMPARE & SWAP

```
AtomicInteger atomic = new AtomicInteger(0);
atomic.getAndIncrement();

boolean applied = false;
do {
    int value = sharedVar.get();
    applied = sharedVar.compareAndSet(
        value, value + 1);
} while ( ! applied );
```

# Concurrent Stack

```
public class ConcurrentStack <E> {
    AtomicReference<Node<E>> top = new AtomicReference<Node<E>>();

    public void push(E item) {
        Node<E> newHead = new Node<E>(item);
        Node<E> oldHead;
        do {
            oldHead = top.get();
            newHead.next = oldHead;
        } while (!top.compareAndSet(oldHead, newHead));
    }

    public E pop() {
        Node<E> oldHead;
        Node<E> newHead;
        do {
            oldHead = top.get();
            if (oldHead == null)
                return null;
            newHead = oldHead.next;
        } while (!top.compareAndSet(oldHead, newHead));
        return oldHead.item;
    }
}
```

# Concurrent Stack

```
public class ConcurrentStack <E> {
    AtomicReference<Node<E>> top = new AtomicReference<Node<E>>();

    public void push(E item) {
        Node<E> newHead = new Node<E>(item);
        Node<E> oldHead;
        do {
            oldHead = top.get();
            newHead.next = oldHead;
        } while (!top.compareAndSet(oldHead, newHead));
    }

    public E pop() {
        Node<E> oldHead;
        Node<E> newHead;
        do {
            oldHead = top.get();
            if (oldHead == null)
                return null;
            newHead = oldHead.next;
        } while (!top.compareAndSet(oldHead, newHead));
        return oldHead.item;
    }
}
```

# Concurrent Stack

```
public class ConcurrentStack <E> {
    AtomicReference<Node<E>> top = new AtomicReference<Node<E>>();

    public void push(E item) {
        Node<E> newHead = new Node<E>(item);
        Node<E> oldHead;
        do {
            oldHead = top.get();
            newHead.next = oldHead;
        } while (!top.compareAndSet(oldHead, newHead));
    }

    public E pop() {
        Node<E> oldHead;
        Node<E> newHead;
        do {
            oldHead = top.get();
            if (oldHead == null)
                return null;
            newHead = oldHead.next;
        } while (!top.compareAndSet(oldHead, newHead));
        return oldHead.item;
    }
}
```

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = factor(i);  
        encodeIntoResponse(resp, factors);  
    }  
}
```

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
    private final AtomicReference<BigInteger> lastNumber  
        = new AtomicReference<BigInteger>();  
    private final AtomicReference<BigInteger[]> lastFactors  
        = new AtomicReference<BigInteger[]>();  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        if (i.equals(lastNumber.get()))  
            encodeIntoResponse(resp, lastFactors.get());  
        else {  
            BigInteger[] factors = factor(i);  
            lastNumber.set(i);  
            lastFactors.set(factors);  
            encodeIntoResponse(resp, factors);  
        }  
    }  
}
```

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
    private final AtomicReference<BigInteger> lastNumber  
        = new AtomicReference<BigInteger>();  
    private final AtomicReference<BigInteger[]> lastFactors  
        = new AtomicReference<BigInteger[]>();  
  
    public void service(ServletRequest req, ServletResponse resp) {  
        BigInteger i = extractFromRequest(req);  
        if (i.equals(lastNumber.get()))  
            encodeIntoResponse(resp, lastFactors.get());  
        else {  
            BigInteger[] factors = factor(i);  
            lastNumber.set(i);  
            lastFactors.set(factors);  
            encodeIntoResponse(resp, factors);  
        }  
    }  
}
```

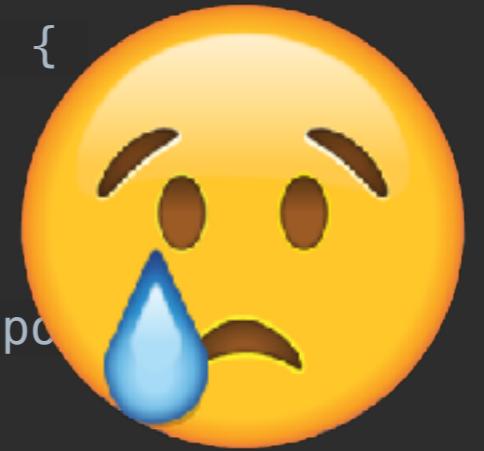


“To preserve state consistency, update related state variables in a **single atomic operation**.”

*–Commandment #1*

# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
    @GuardedBy("this") private BigInteger lastNumber;  
    @GuardedBy("this") private BigInteger[] lastFactors  
        = new BigInteger[]{};  
  
    public synchronized void service(ServletRequest req, ServletResponse resp)  
    {  
        BigInteger i = extractFromRequest(req);  
        if (i == lastNumber)  
            encodeIntoResponse(resp, lastFactors);  
        else {  
            BigInteger[] factors = factor(i);  
            lastNumber = i;  
            lastFactors = factors;  
            encodeIntoResponse(resp, factors);  
        }  
    }  
}
```



# Factorizer

```
public class Factorizer extends GenericServlet implements Servlet {  
    @GuardedBy("this") private BigInteger lastNumber;  
    @GuardedBy("this") private BigInteger[] lastFactors  
        = new BigInteger[]{};  
  
    public synchronized void service(ServletRequest req, ServletResponse resp)  
        throws ServletException, IOException {  
        BigInteger i = extractFromRequest(req);  
        BigInteger[] factors = null;  
        synchronized (this) {  
            if (i.equals(lastNumber)) {  
                factors = lastFactors.clone();  
            }  
        }  
        if (factors == null) {  
            factors = factor(i);  
            synchronized (this) {  
                lastNumber = i;  
                lastFactors = factors.clone();  
            }  
        }  
        encodeIntoResponse(resp, factors);  
    }  
}
```



# Synchronized

```
public class SynchronizedDemo {  
    public synchronized void update() { }  
}
```

# Synchronized

```
public class SynchronizedDemo {  
    public synchronized void update() {  
        synchronized (this) {  
            }  
    }  
}
```

# Synchronized

```
public class SynchronizedDemo {  
    private Object lock = new Object();  
    public synchronized void update() {  
        synchronized (lock) {  
            }  
    }  
}
```

# Synchronized

```
public class SynchronizedDemo {  
    @GuardedBy("lock") Integer counter1;  
    private Object lock = new Object();  
    public void update() {  
        synchronized (lock) {  
            if (counter1 > Integer.MAX_VALUE) {  
                return;  
            }  
            counter1 = counter1 + 1;  
        }  
    }  
}
```

# Synchronized

```
public class SynchronizedDemo {  
    @GuardedBy("lock") Integer counter1;  
    @GuardedBy("lock2") Integer counter2;  
    private Object lock = new Object();  
    private Object lock2 = new Object();  
    public void update() {  
    }  
    public void update2() {  
        synchronized (lock2) {  
            if (counter2 > Integer.MAX_VALUE) {  
                return;  
            }  
            counter2 = counter2 + 1;  
        }  
    }  
}
```

# Synchronized

```
public class SynchronizedDemo extends BaseUpdate {  
    public synchronized void update() {  
        super.update();  
    }  
}  
  
class BaseUpdate {  
    public synchronized void update() {}  
}
```

# Locking

# Locking

- Prefer explicit locking. Exception - superclass uses synchronised keyword, private lock will be unusable for subclasses

# Locking

- Prefer explicit locking. Exception - superclass uses synchronised keyword, private lock will be unusable for subclasses
- Locks are reentrant.

# Locking

- Prefer explicit locking. Exception - superclass uses synchronised keyword, private lock will be unusable for subclasses
- Locks are reentrant.
- Add annotation @GuardedBy("CorrectLock") for documentation.

# Locking

- Prefer explicit locking. Exception - superclass uses synchronised keyword, private lock will be unusable for subclasses
- Locks are reentrant.
- Add annotation `@GuardedBy("CorrectLock")` for documentation.
- Perform all operations on a single variable using the same lock.

# Locking

- Prefer explicit locking. Exception - superclass uses synchronised keyword, private lock will be unusable for subclasses
- Locks are reentrant.
- Add annotation `@GuardedBy("CorrectLock")` for documentation.
- Perform all operations on a single variable using the same lock.
- Perform all operations on related variables using the same lock. (i.e. Unrelated variables can have different locks)

# LOCK COARSENING

```
synchronized ( buffer ) {  
    buffer.add(x);  
}  
foo = bar;  
synchronized ( buffer ) {  
    buffer.add(y);  
}
```



```
synchronized ( buffer ) {  
    buffer.add(x);  
    foo = bar;  
    buffer.add(y);  
}
```

# Thread Confinement

```
BehaviorSubject<String> timerSubject =  
BehaviorSubject.create();  
  
Observable  
    .timer(4, TimeUnit.SECONDS)  
    .subscribe(aLong -> {  
        timerSubject.onNext(String.valueOf(aLong));  
    });
```

# Thread Confinement

```
BehaviorSubject<String> timerSubject =  
BehaviorSubject.create();
```

```
Observable  
    .timer(4, TimeUnit.SECONDS, Schedulers.computation())  
    .subscribe(aLong -> {  
        timerSubject.onNext(String.valueOf(aLong));  
    });
```

\* *All methods except the onSubscribe, onNext, onError and onComplete are thread-safe.*

\* *Use {@link #toSerialized()} to make these methods thread-safe as well.*



# Thread Confinement



```
Subject<String> timerSubject =  
BehaviorSubject.create().serialized();
```

```
Observable  
    .timer(4, TimeUnit.SECONDS, Schedulers.computation())  
    .subscribe(aLong -> {  
        timerSubject.onNext(String.valueOf(aLong));  
    });
```

\* *All methods except the `onSubscribe`, `onNext`, `onError` and `onComplete` are thread-safe.*

\* *Use `{@link #toSerialized()}` to make these methods thread-safe as well.*

# Thread Confinement

```
Scheduler scheduler = Schedulers.single();
BehaviorSubject<String> timerSubject
= BehaviorSubject.create();
```



```
Observable
    .timer(4, TimeUnit.SECONDS, Schedulers.computation())
    .observeOn(scheduler)
    .subscribe(aLong -> {
        timerSubject.onNext(String.valueOf(aLong));
    });
}
```

\* <p>All methods except the `onSubscribe`, `onNext`, `onError` and `onComplete` are thread-safe.

\* Use `{@link #toSerialized()}` to make these methods thread-safe as well.

# RX-Java Thread Safety

# RX-Java Thread Safety

- Subjects are not thread safe.

# RX-Java Thread Safety

- Subjects are not thread safe.
- Most of the operators with single observable are not thread safe ex - take(N)

# RX-Java Thread Safety

- Subjects are not thread safe.
- Most of the operators with single observable are not thread safe ex - take(N)
- Most of the operators with multiple observable are thread safe ex- merge, concat

# RX-Java Thread Safety

- Subjects are not thread safe.
- Most of the operators with single observable are not thread safe ex - take(N)
- Most of the operators with multiple observable are thread safe ex- merge, concat
- Details - [blog](#)

# Visibility

```
boolean run = false;
int x = 42;

public void thread1() {
    Thread thread = new Thread(new Runnable() {
        @Override public void run() {
            while (!run) { }
            System.out.println(x);
        }
    });
    thread.start();
}

public void thread2() {
    x = 100;
    run = true;
}
```

# Visibility



```
boolean run = false;
int x = 42;

public void thread1() {
    Thread thread = new Thread(new Runnable() {
        @Override public void run() {
            while (!run) { }
            System.out.println(x);
        }
    });
    thread.start();
}

public void thread2() {
    x = 100;
    run = true;
}
// Output - Never terminate, 42, 100
```

# Visibility



```
volatile boolean run = false;
volatile int x = 42;

public void thread1() {
    Thread thread = new Thread(new Runnable() {
        @Override public void run() {
            while (!run) { }
            System.out.println(x);
        }
    });
    thread.start();
}

public void thread2() {
    x = 100;
    run = true;
}
// Output - 100
```

# WAIT & NOTIFY

## PRODUCER

```
synchronized (lock) {  
    sharedData = ...;  
    sharedDone = true;  
    lock.notify();  
}
```

## CONSUMER

```
synchronized (lock) {  
    while ( !sharedDone ) {  
        lock.wait();  
    }  
    print(sharedData);  
}
```

# DOESN'T PRESERVE ORDER

## PRODUCER

```
synchronized (lock) {  
    sharedData = ...;  
    sharedDone = true;  Reorder!  
    lock.notify();  
}  
  
↓  
  
synchronized (lock) {  
    sharedDone = true;  
    sharedData = ...;  
    lock.notify();  
}
```

## CONSUMER

```
synchronized (lock) {  
    while ( !sharedDone ) {  
        lock.wait();  
    }  
    print(sharedData);  
}
```

# SPURIOUS NOTIFICATIONS

## \*INCORRECT\* CONSUMER

```
synchronized (lock) {  
    if ( !sharedDone ) {  
        lock.wait();  
    }  
    print(sharedData);  
}
```

## \*CORRECT\* CONSUMER

```
synchronized (lock) {  
    while ( !sharedDone ) {  
        lock.wait();  
    }  
    print(sharedData);  
}
```

“Locking can guarantee both visibility and atomicity; volatile variables can only guarantee visibility.”

*–Commandment #2*

# Put if Absent

```
public class MapExample<K, V> {  
    Map<K, V> myMap = new ConcurrentHashMap<>();  
  
    public V addIfAbsent(K key, V value) {  
        if (!myMap.containsKey(key)) {  
            myMap.put(key, value);  
        }  
        return myMap.get(key);  
    }  
}
```

# Put if Absent

```
public class MapExample<K, V> {  
  
    Map<K, V> myMap = new ConcurrentHashMap<>();  
  
    public V addIfAbsent(K key, V value) {  
        if (!myMap.containsKey(key)) {  
            myMap.put(key, value);  
        } // Non - Atomic  
        return myMap.get(key);  
    }  
}
```



# Put if Absent

```
public class MapExample<K, V> {  
    Map<K, V> myMap = new ConcurrentHashMap<>();  
  
    public synchronized V addIfAbsent(K key, V value) {  
        if (!myMap.containsKey(key)) {  
            myMap.put(key, value);  
        }  
        return myMap.get(key);  
    }  
}
```

# Put if Absent

```
public class MapExample<K, V> {  
    Map<K, V> myMap = new ConcurrentHashMap<>();  
  
    public synchronized V addIfAbsent(K key, V value) {  
        if (!myMap.containsKey(key)) {  
            myMap.put(key, value);  
        }  
        return myMap.get(key);  
    }  
}  
  
// Concurrent hash map is synchronised using another lock.
```



# Put if Absent

```
public class MapExample<K, V> {  
    Map<K, V> myMap = new ConcurrentHashMap<>();  
  
    public synchronized V addIfAbsent(K key, V value) {  
        synchronized (myMap) {  
            if (!myMap.containsKey(key)) {  
                myMap.put(key, value);  
            }  
            return myMap.get(key);  
        }  
    }  
}  
  
// Internal lock change and what if the lock is private
```



# Put if Absent

```
public class MapExample<K, V> {  
    MyMap<K, V> myMap = new MyMap<>(new HashMap<>());  
}  
  
class MyMap<K, V> implements Map<K, V> {  
  
    @GuardedBy("lock") private Map<K, V> delegate;  
    private Object lock = new Object();  
  
    public MyMap(Map<K, V> map) {  
        delegate = map;  
    }  
  
    @Override public V put(K key, V value) {  
        synchronized (lock) {  
            delegate.put(key, value);  
        }  
    }  
  
    public V addIfAbsent(K key, V value) {  
        synchronized (lock) {  
            if (!delegate.containsKey(key)) {  
                delegate.put(key, value);  
            }  
            return delegate.get(key);  
        }  
    }  
}
```



# Put if Absent

```
public class MapExample<K, V> {  
    MyMap<K, V> myMap = new MyMap<>(new HashMap<>());  
}  
  
class MyMap<K, V> implements Map<K, V> {  
  
    @GuardedBy("lock") private Map<K, V> delegate;  
    private Object lock = new Object();  
  
    public MyMap(Map<K, V> map) {  
        delegate = map;  
    }  
  
    @Override public V put(K key, V value) {  
        synchronized (lock) {  
            delegate.put(key, value);  
        }  
    }  
  
    public V addIfAbsent(K key, V value) {  
        synchronized (lock) {  
            if (!delegate.containsKey(key)) {  
                delegate.put(key, value);  
            }  
            return delegate.get(key);  
        }  
    }  
}
```



# Put if Absent



```
public class MapExample<K, V> {  
    MyMap<K, V> myMap = new MyMap<>(new HashMap<>());  
}  
  
class MyMap<K, V> implements Map<K, V> {  
  
    @GuardedBy("lock") private Map<K, V> delegate;  
    private Object lock = new Object();  
  
    public MyMap(Map<K, V> map) {  
        delegate = map;  
    }  
  
    @Override public V put(K key, V value) {  
        synchronized (lock) {  
            delegate.put(key, value);  
        }  
    }  
  
    public V addIfAbsent(K key, V value) {  
        synchronized (lock) {  
            if (!delegate.containsKey(key)) {  
                delegate.put(key, value);  
            }  
            return delegate.get(key);  
        }  
    }  
}
```

# Put if Absent

```
public class MapExample<K, V> {  
    MyMap<K, V> myMap = new MyMap<>(new HashMap<>());  
}  
  
class MyMap<K, V> implements Map<K, V> {  
  
    @GuardedBy("lock") private Map<K, V> delegate;  
    private Object lock = new Object();  
  
    public MyMap(Map<K, V> map) {  
        delegate = map;  
    }  
  
    @Override public V put(K key, V value) {  
        synchronized (lock) {  
            delegate.put(key, value);  
        }  
    }  
  
    public V addIfAbsent(K key, V value) {  
        synchronized (lock) {  
            if (!delegate.containsKey(key)) {  
                delegate.put(key, value);  
            }  
            return delegate.get(key);  
        }  
    }  
}
```



# Put if Absent

```
public class MapExample<K, V> {
    MyMap<K, V> myMap = new MyMap<>(new HashMap<>());
}

class MyMap<K, V> implements Map<K, V> {

    @GuardedBy("lock") private Map<K, V> delegate;
    private Object lock = new Object();

    public MyMap(Map<K, V> map) {
        delegate = map;
    }

    @Override public V put(K key, V value) {
        synchronized (lock) {
            delegate.put(key, value);
        }
    }

    public V addIfAbsent(K key, V value) {
        synchronized (lock) {
            if (!delegate.containsKey(key)) {
                delegate.put(key, value);
            }
            return delegate.get(key);
        }
    }
}
```



# Publishing

```
public class Publishing {
    private HashMap<String, String> mutableMap = new HashMap<()>;
    private ImmutableHolder<String, String> immutableHolder = new ImmutableHolder<mutableMap>();

    @NotThreadSafe
    public void update() {
        immutableHolder.getStringMap().put("junk", "Junk");
    }
}

class ImmutableHolder<K, V> {
    private final Map<K, V> internalMap;

    ImmutableHolder(Map<K, V> internalMap) {
        this.internalMap = internalMap;
    }

    public Map<K, V> getStringMap() {
        return internalMap;
    }

    public synchronized void putValue(K key, V value) {
        internalMap.put(key, value);
    }
}
```

# Publishing

```
public class Publishing {  
  
    private ImmutableHolder<String, String> immutableHolder;  
  
    @NotThreadSafe  
    public void update() {  
        immutableHolder.getStringMap().put("junk", "junk");  
    }  
  
    class ImmutableHolder<K, V> {  
        private final Map<K, V> internalMap;  
  
        ImmutableHolder(Map<K, V> internalMap) {  
            this.internalMap = internalMap;  
        }  
  
        public Map<K, V> getStringMap() {  
            return internalMap;  
        }  
  
        public synchronized void putValue(K key, V value) {  
            internalMap.put(key, value);  
        }  
    }  
}
```



# Publishing

```
public class Publishing {  
  
    private final Map<String, String> internalMap = new HashMap<String, String>();  
  
    private ImmutableHolder<String, String> immutableHolder;  
  
    @NotThreadSafe  
    public void update() {  
        immutableHolder.getStringMap().put("junk", "junk");  
    }  
  
    class ImmutableHolder<K, V> {  
        private final Map<K, V> internalMap;  
  
        ImmutableHolder(Map<K, V> internalMap) {  
            this.internalMap = internalMap;  
        }  
  
        public Map<K, V> getStringMap() {  
            return internalMap.clone();  
        }  
  
        public synchronized void putValue(K key, V value) {  
            internalMap.put(key, value);  
        }  
    }  
}
```



# Publishing

```
public class Publishing {

    private HashMap<String, String> mutableMap = new HashMap<>();
    private ImmutableHolder<String, String> immutableHolder = new ImmutableHolder(mutableMap);

    @NotThreadSafe
    public void update() {
        mutableMap.put("junk", "junk");
    }

    class ImmutableHolder<K, V> {
        private final Map<K, V> internalMap;

        ImmutableHolder(Map<K, V> internalMap) {
            this.internalMap = internalMap;
        }

        public Map<K, V> getStringMap() {
            return internalMap.clone();
        }

        public synchronized void putValue(K key, V value) {
            internalMap.put(key, value);
        }
    }
}
```



# Publishing

```
public class Publishing {

    private HashMap<String, String> mutableMap = new HashMap<>();
    private ImmutableHolder<String, String> immutableHolder = new ImmutableHolder(mutableMap);

    @NotThreadSafe
    public void update() {
        mutableMap.put("junk", "junk");
    }

    class ImmutableHolder<K, V> {
        private final Map<K, V> internalMap;

        ImmutableHolder(Map<K, V> internalMap) {
            this.internalMap = internalMap.clone();
        }

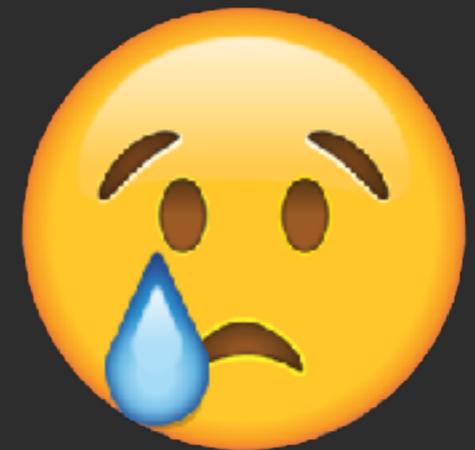
        public Map<K, V> getStringMap() {
            return internalMap.clone();
        }

        public synchronized void putValue(K key, V value) {
            internalMap.put(key, value);
        }
    }
}
```



# Publishing

```
public class Publishing {  
  
    private HashMap<String, String> mutableMap = new HashMap<>();  
    private ImmutableHolder<String, String> immutableHolder = new  
    ImmutableHolder(mutableMap);  
  
    @NotThreadSafe  
    public void update() {  
        mutableMap.put("junk", "junk"); // no effect  
        immutableHolder.getStringMap().put("junk", "junk"); // Crash.  
    }  
  
    class ImmutableHolder<K, V> {  
        private final Map<K, V> internalMap;  
  
        ImmutableHolder(Map<K, V> internalMap) {  
            this.internalMap = Collections.unmodifiableMap(internalMap);  
        }  
  
        public Map<K, V> getStringMap() {  
            return internalMap.clone();  
        }  
    }  
}
```



# Publishing

# Publishing

- Immutable objects can be published anywhere.

# Publishing

- Immutable objects can be published anywhere.
- Effectively immutable objects (mutable objects which we will not change) can be published using “safe publishing” - volatile, final, static initializer.

# Publishing

- Immutable objects can be published anywhere.
- Effectively immutable objects (mutable objects which we will not change) can be published using “safe publishing” - volatile, final, static initializer.
- Mutable variables must be synchronized

# Annotations

- `@ThreadSafe`, `@NotThreadSafe`
- `@Immutable`
- `@GuardedBy`
  - `@GuardedBy("this")`
  - `@GuardedBy("fieldName")`
  - `@GuardedBy("methodName()")`
- `@MainThread`, `@WorkerThread`

# Summary

# Summary

- Encapsulation is your best friend.

# Summary

- Encapsulation is your best friend.
- Document thread safety.

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.
- Don't let constructor escape this

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.
- Don't let constructor escape this
- Publication:

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.
- Don't let constructor escape this
- Publication:
  - Immutable objects can be published anywhere

# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.
- Don't let constructor escape this
- Publication:
  - Immutable objects can be published anywhere
  - Effectively immutable objects can be safely published.

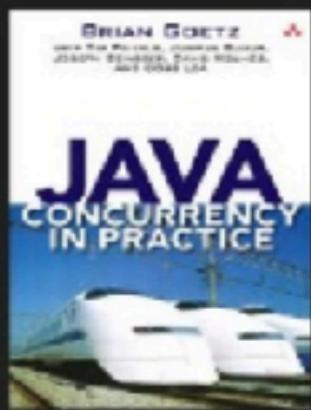
# Summary

- Encapsulation is your best friend.
- Document thread safety.
- Use composition to extend a synchronised class.
- Use immutable variables.
- Don't let constructor escape this
- Publication:
  - Immutable objects can be published anywhere
  - Effectively immutable objects can be safely published.
  - Mutable object needs to be kept behind a lock.

# TODO

- Java concurrency in Practice - Chapter 2-5, Chapter 13, 15, and 16.
- Devoxx talk on [concurrency](#).
- Write code yourself?
- How to test concurrency? - [JCStress](#)

# REFERENCES



## JAVA CURRENCY IN PRACTICE

Brian Goetz et al

## JAVA MEMORY MODEL PRAGMATICS

<http://virtualjug.com/java-memory-model-pragmatics/>

Aleksey Shipilëv



## CLOSE ENCOUNTERS OF THE JMM KIND

<https://shipilev.net/blog/2016/close-encounters-of-jmm-kind/>

Aleksey Shipilëv

# Questions??